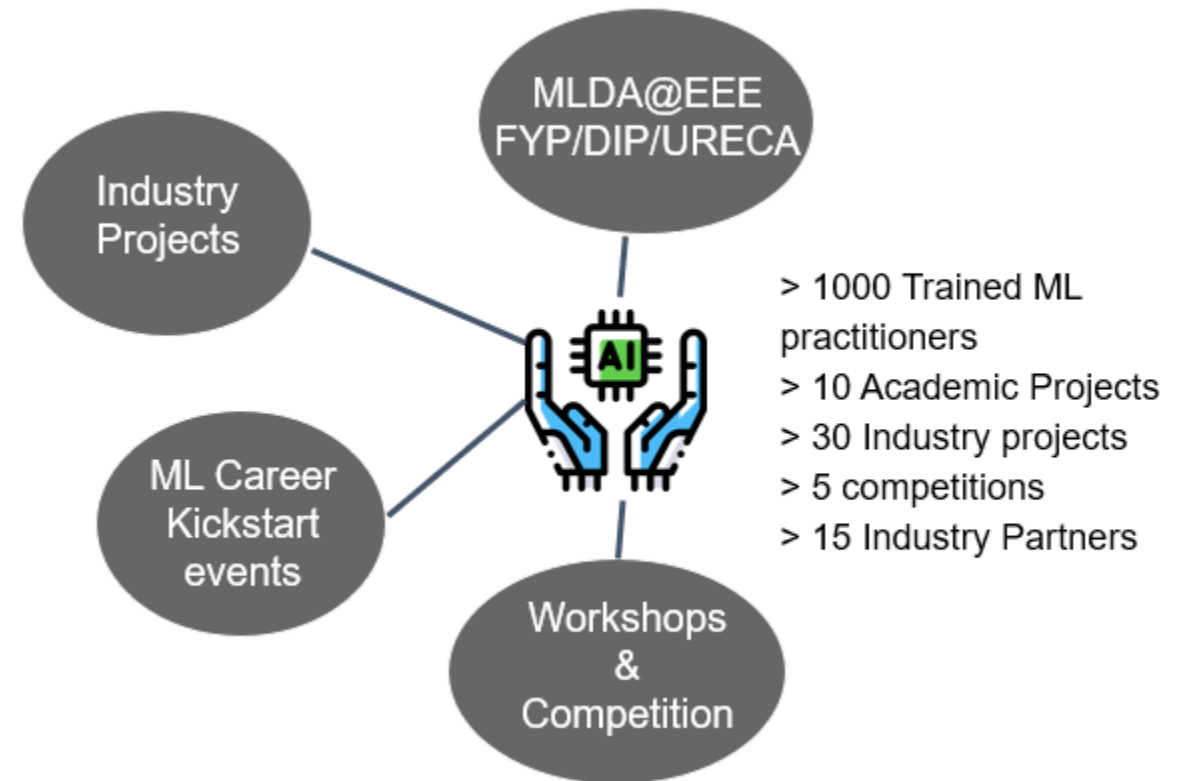
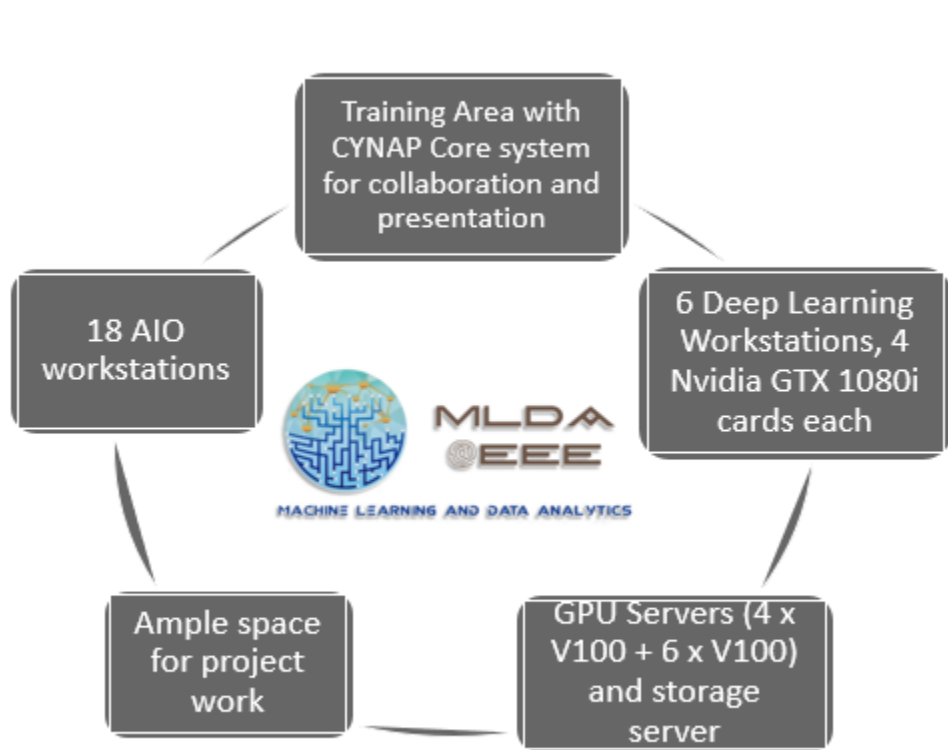


Generative Adversarial Networks (GANs) Workshop

**Presented by:
Chan Yi Xuan
EEE, Y3**

MLDA's Mission

Provide an integrated platform for EEE/IEM students to learn and implement Machine Learning, Data Science & AI, as well as facilitate connections with the industry.



GENERATIVE ADVERSARIAL NETWORKS (GANs)


Which one is FAKE?



GENERATIVE ADVERSARIAL NETWORKS (GANs)

Generative Models

Given training data, generative models aim at **learning the true data distribution** of the training set to **generate new data points** from this distribution with **some variations**.



Training samples

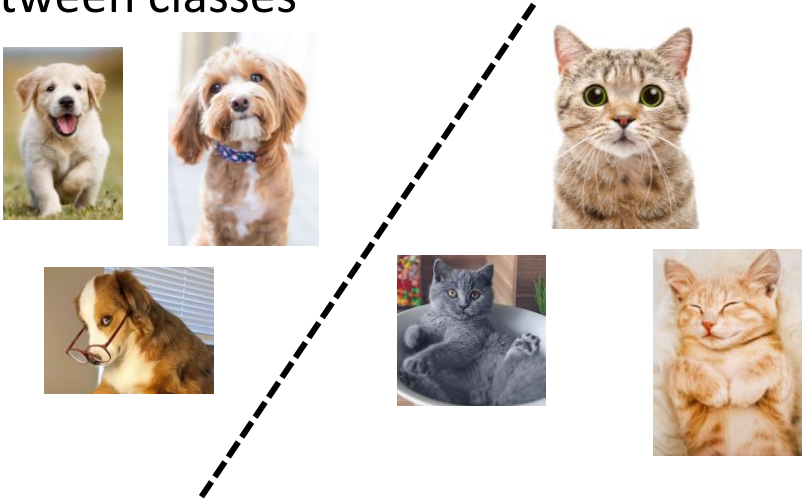


Generated samples

Classifier Model vs. Generative Model

Classifier Model

- Goal: Models the **decision boundary** between classes



Generative Model

- Goal: Models the **actual distribution** of each classes



Generative Models

Given training data, generative models aim at **learning the true data distribution of the training set to **generate new data points** from this distribution with **some variations**.**

Two famous deep generative model algorithms:

- **Variational Autoencoder (VAE)**
- **Generative Adversarial Network (GAN)**

Generative Models

Given training data, generative models aim at **learning the true data distribution of the training set to **generate new data points** from this distribution with **some variations**.**

Two famous deep generative model algorithms:

- **Variational Autoencoder (VAE)**
- **Generative Adversarial Network (GAN)**
 - **Applications of GANs**
 - **Concept and theory behind GANs**
 - **Properties of DCGAN**

GANs: Cool Applications



Face Generation StyleGAN2

More from
<https://thispersondoesnotexist.com/>

GANs: Cool Applications



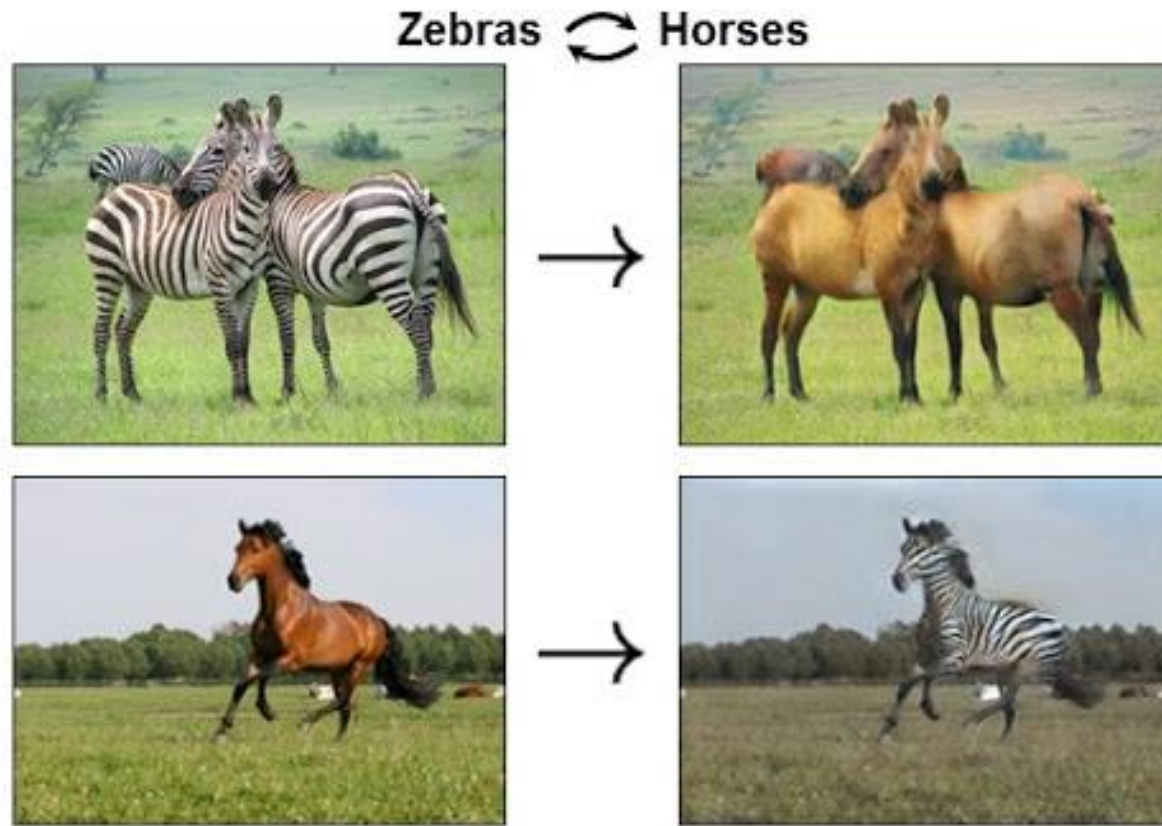
**Cat Generation
StyleGAN2**

GANs: Cool Applications



**Super Resolution
SRGAN**

GANs: Cool Applications



**Image Translation
CycleGAN**

Generative Adversarial Networks

Generator

Learns to make **fakes** that
look **real**

Discriminator

Learns to distinguish **real**
from **fake**


Generative Adversarial Networks

Generator

vs.

Discriminator

Learns to make **fakes** that
look **real**



**Art
Forger**

vs.

**Art
Critic**



GENERATIVE ADVERSARIAL NETWORKS (GANs)

Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples


Output: Probability whether the sample is real

Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real




Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real




Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real




Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real




Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real



Discriminator

Discriminator: A classifier which can distinguish between different classes

Input: Real or Fake Samples

Output: Probability whether the sample is real

$$P(\underset{\text{Class}}{Y} \mid \underset{\text{Features}}{X})$$

Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise


Output: Fake data

Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data




Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Noise Vector

Noise Vector

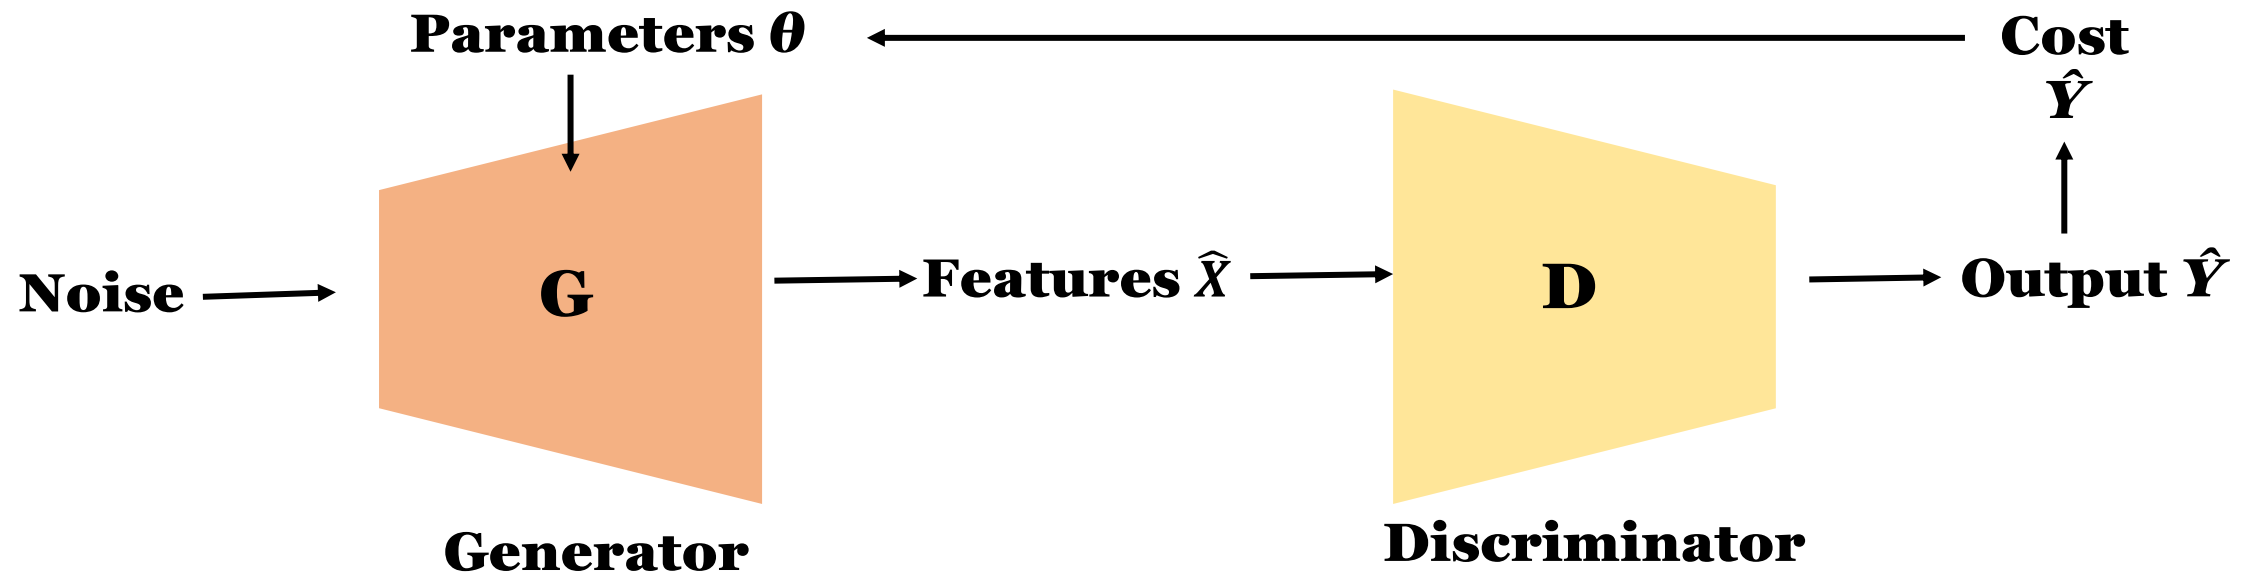
- Also known as **latent variable**, are those variables that are important for a domain but are not directly observable.
- Random noise vector's dimensionality is **smaller** than the target output's dimensionality
- Noise helps GAN to produce a wide variety of data

Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

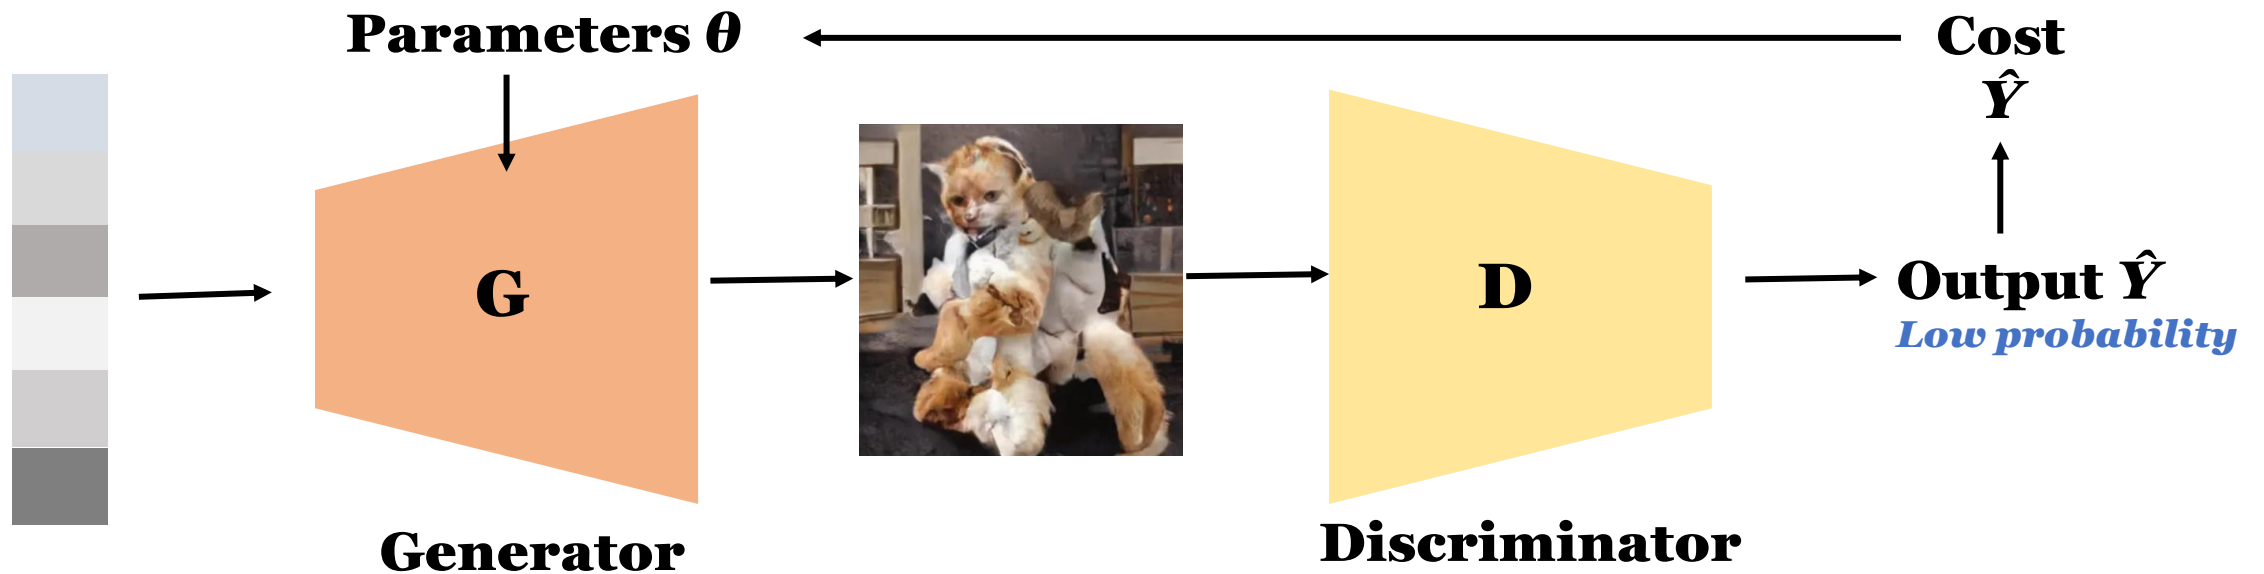


Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

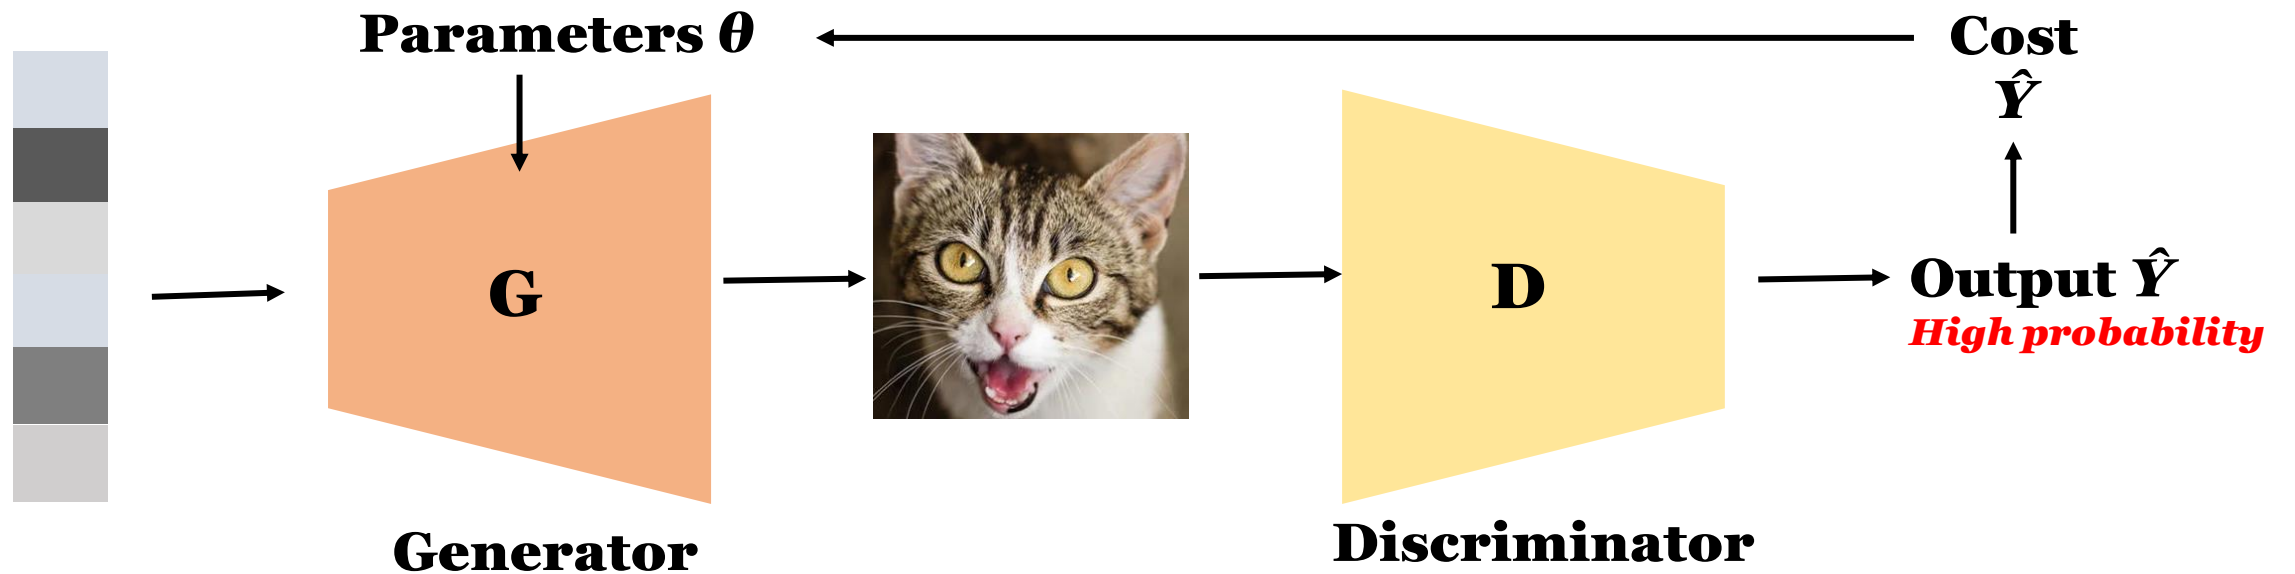


Generator

Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data



Generator


Generator: Model that is used to generate new plausible examples from the problem domain

Input: Random noise

Output: Fake data

$P(X | Y)$
Features **Class**

**Approximate the
distribution of
problem domain**



<https://poloclub.github.io/ganlab/>

Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

m: sample size
y: label
 \hat{y} : prediction

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

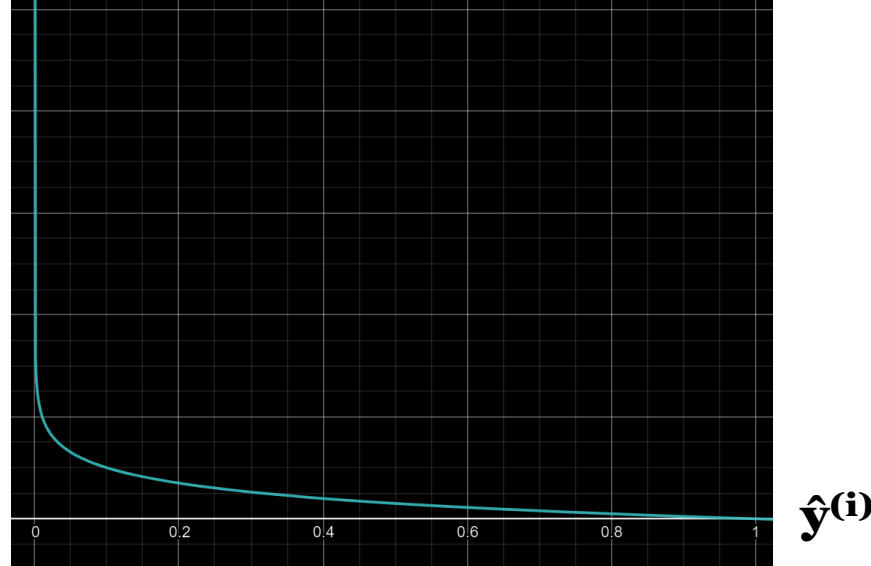
m: sample size
y: label
 \hat{y} : prediction

$$BCE = -\frac{1}{m} \sum_{i=1}^m \boxed{y^{(i)} \log \hat{y}^{(i)}} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Label y = 1 (real)

Note: $\log(1) = 0$

$\log \hat{y}^{(i)}$



Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss


m: sample size
y: label
 \hat{y} : prediction

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + \boxed{(1 - y^{(i)}) \log(1 - \hat{y}^{(i)})}$$

Label y = 0 (fake)

Note: $\log(1) = 0$

$\log(1 - \hat{y}^{(i)})$



Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

BCE Loss -> Minimax Loss

$$\min_d \max_g - \left[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z)))) \right]$$



Discriminator:
Minimize cost

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Generator:
Maximize cost

Cost Function

Loss Function: Binary Cross Entropy (BCE) Loss

$$BCE = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

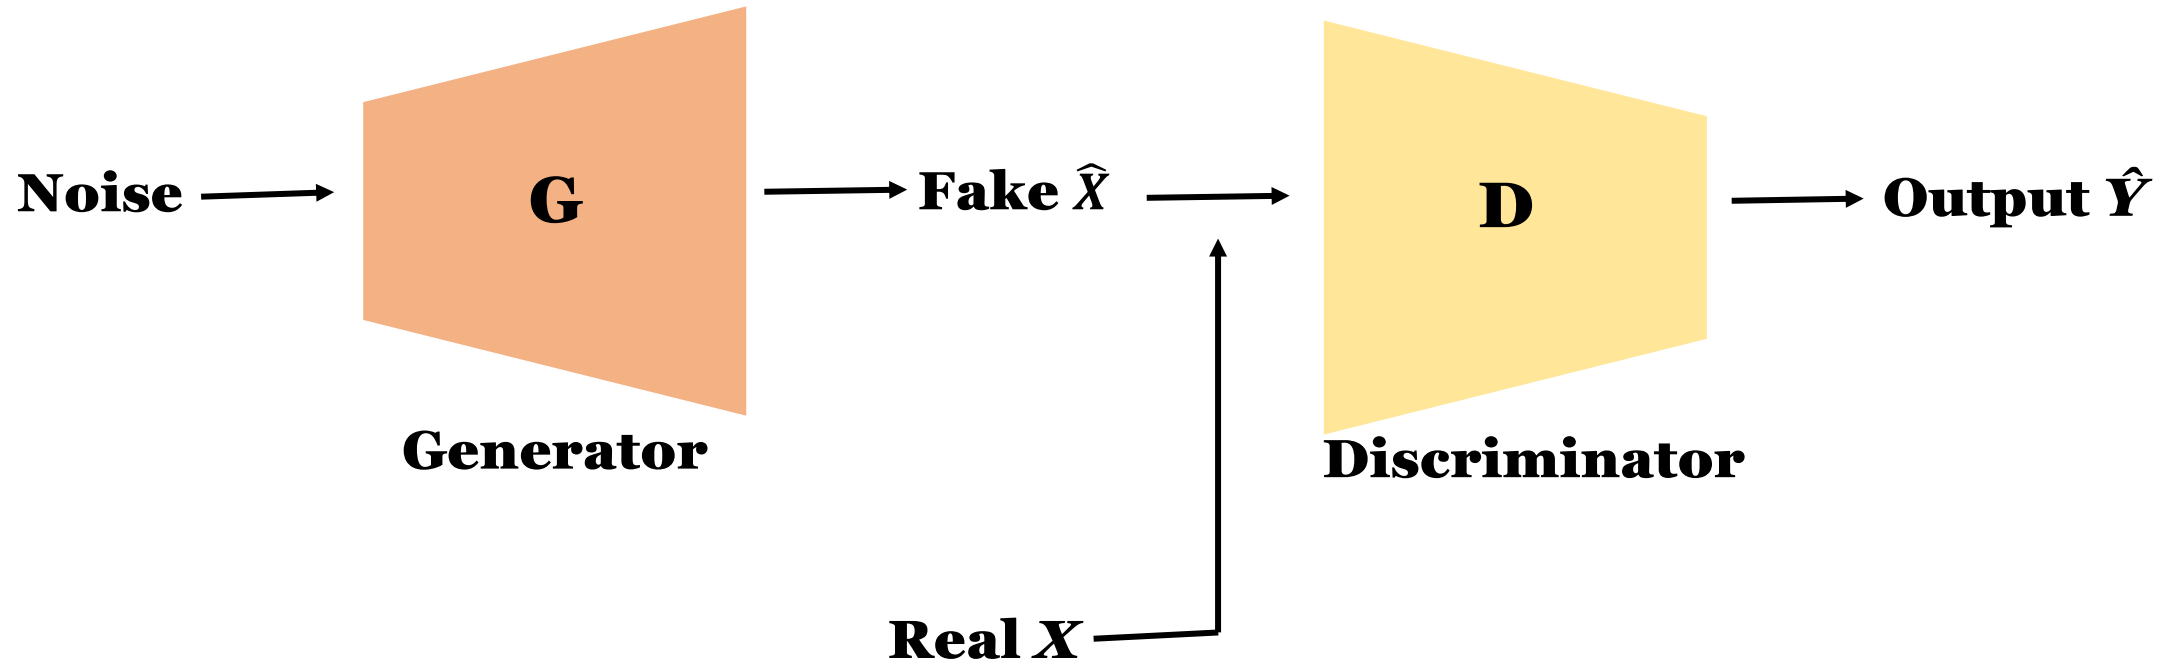
BCE Loss -> Minimax Loss

$$\min_d \max_g - \left[\mathbb{E}(\log(d(x))) + \mathbb{E} \left(1 - \log(d(g(z))) \right) \right]$$

OR...

$$\min_g \max_d \left[\mathbb{E}(\log(d(x))) + \mathbb{E} \left(1 - \log(d(g(z))) \right) \right]$$

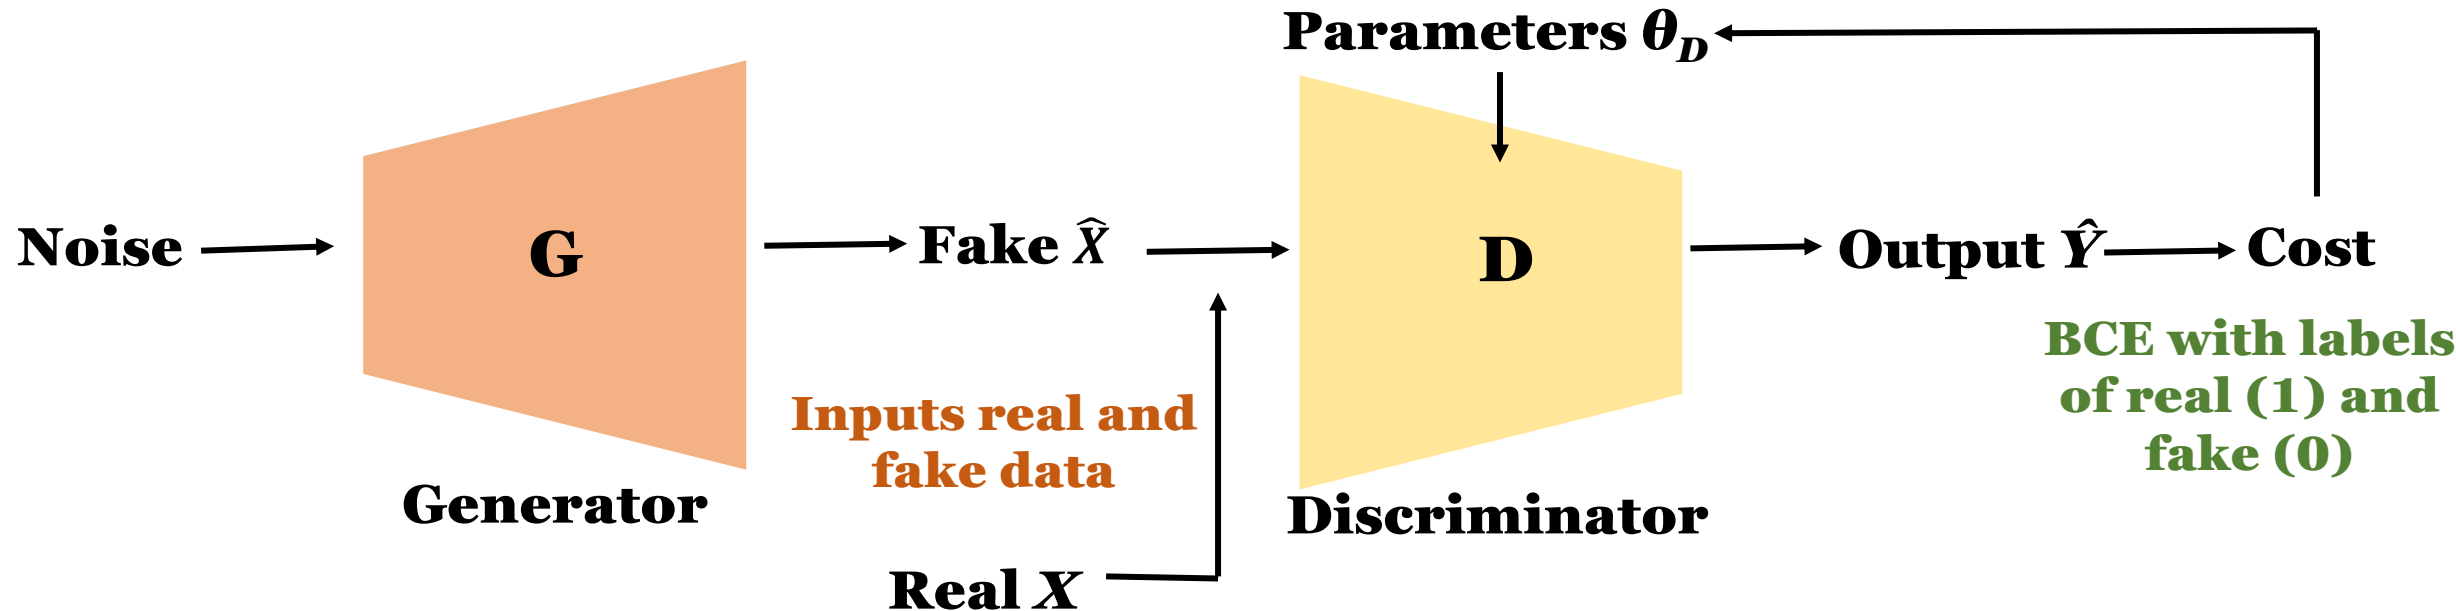
Complete GAN Model



Complete GAN Model

Training Discriminator

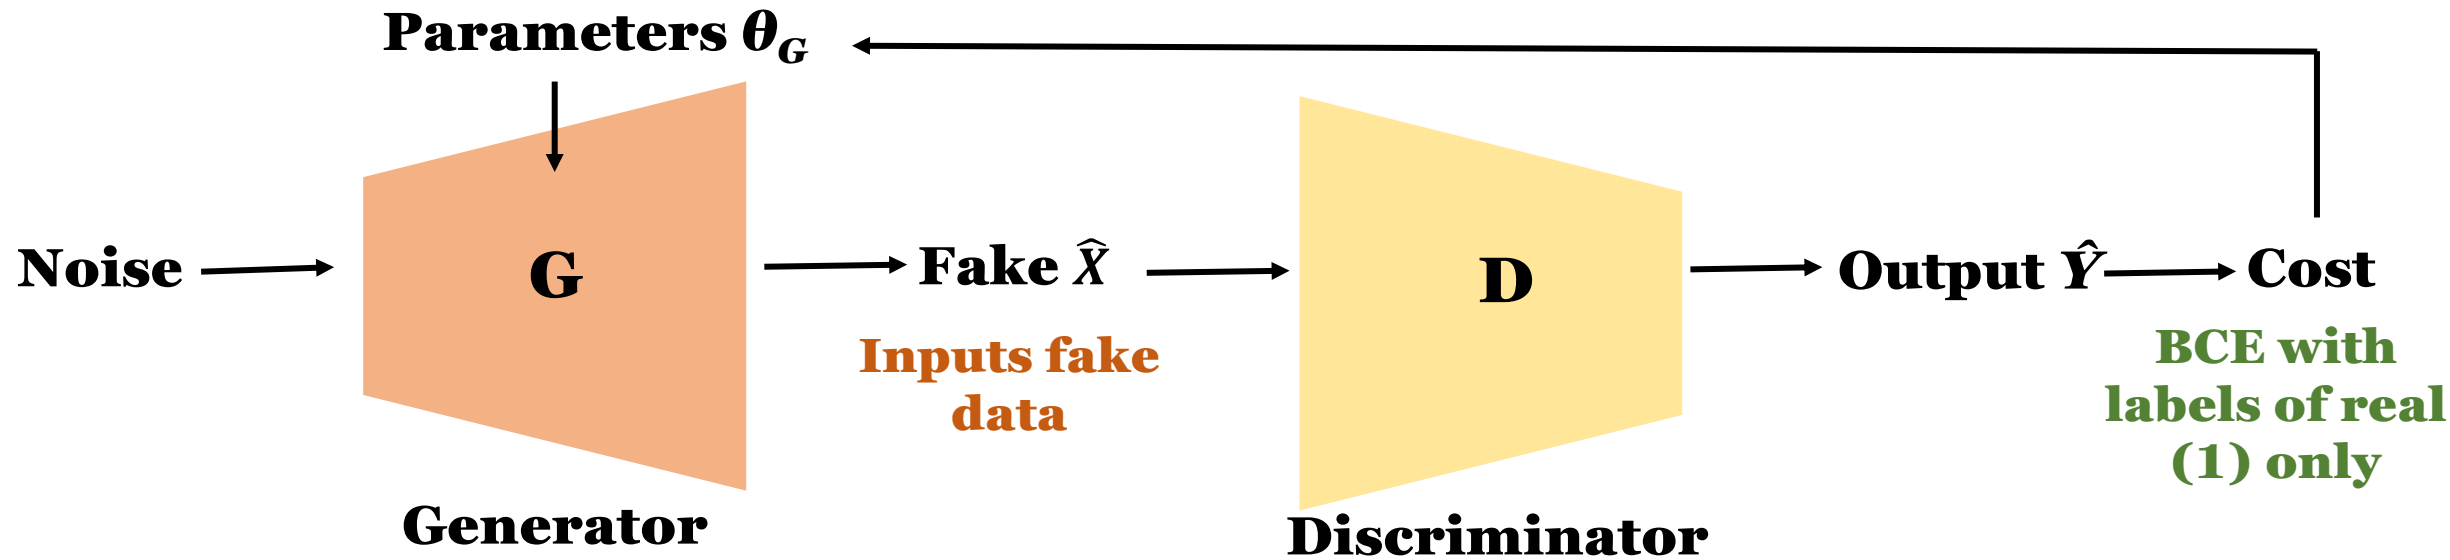
Keep generator constant



Complete GAN Model

Training Generator

Keep discriminator constant



GAN is trained in an **alternative fashion**


GAN is trained in an **alternative fashion**

Do not use discriminator that is too strong!




DCGAN

Generator



Discriminator




- Replace all pooling layers with **convolutional stride**
- Use **transposed convolution** for upsampling
- Use **batch normalization** in both the generator and the discriminator
- Remove fully connected hidden layers
- Use **ReLU activation** in generator for all layers except for the output, which uses **Tanh**
- Use **LeakyReLU** activation in the discriminator for all layers

Credit: <https://towardsdatascience.com/gan-by-example-using-keras-on-tensorflow-backend-1a6d515a60d0>


DCGAN

Discriminator: Convolution ~ downsampling




DCGAN

Discriminator: Convolution ~ downsampling



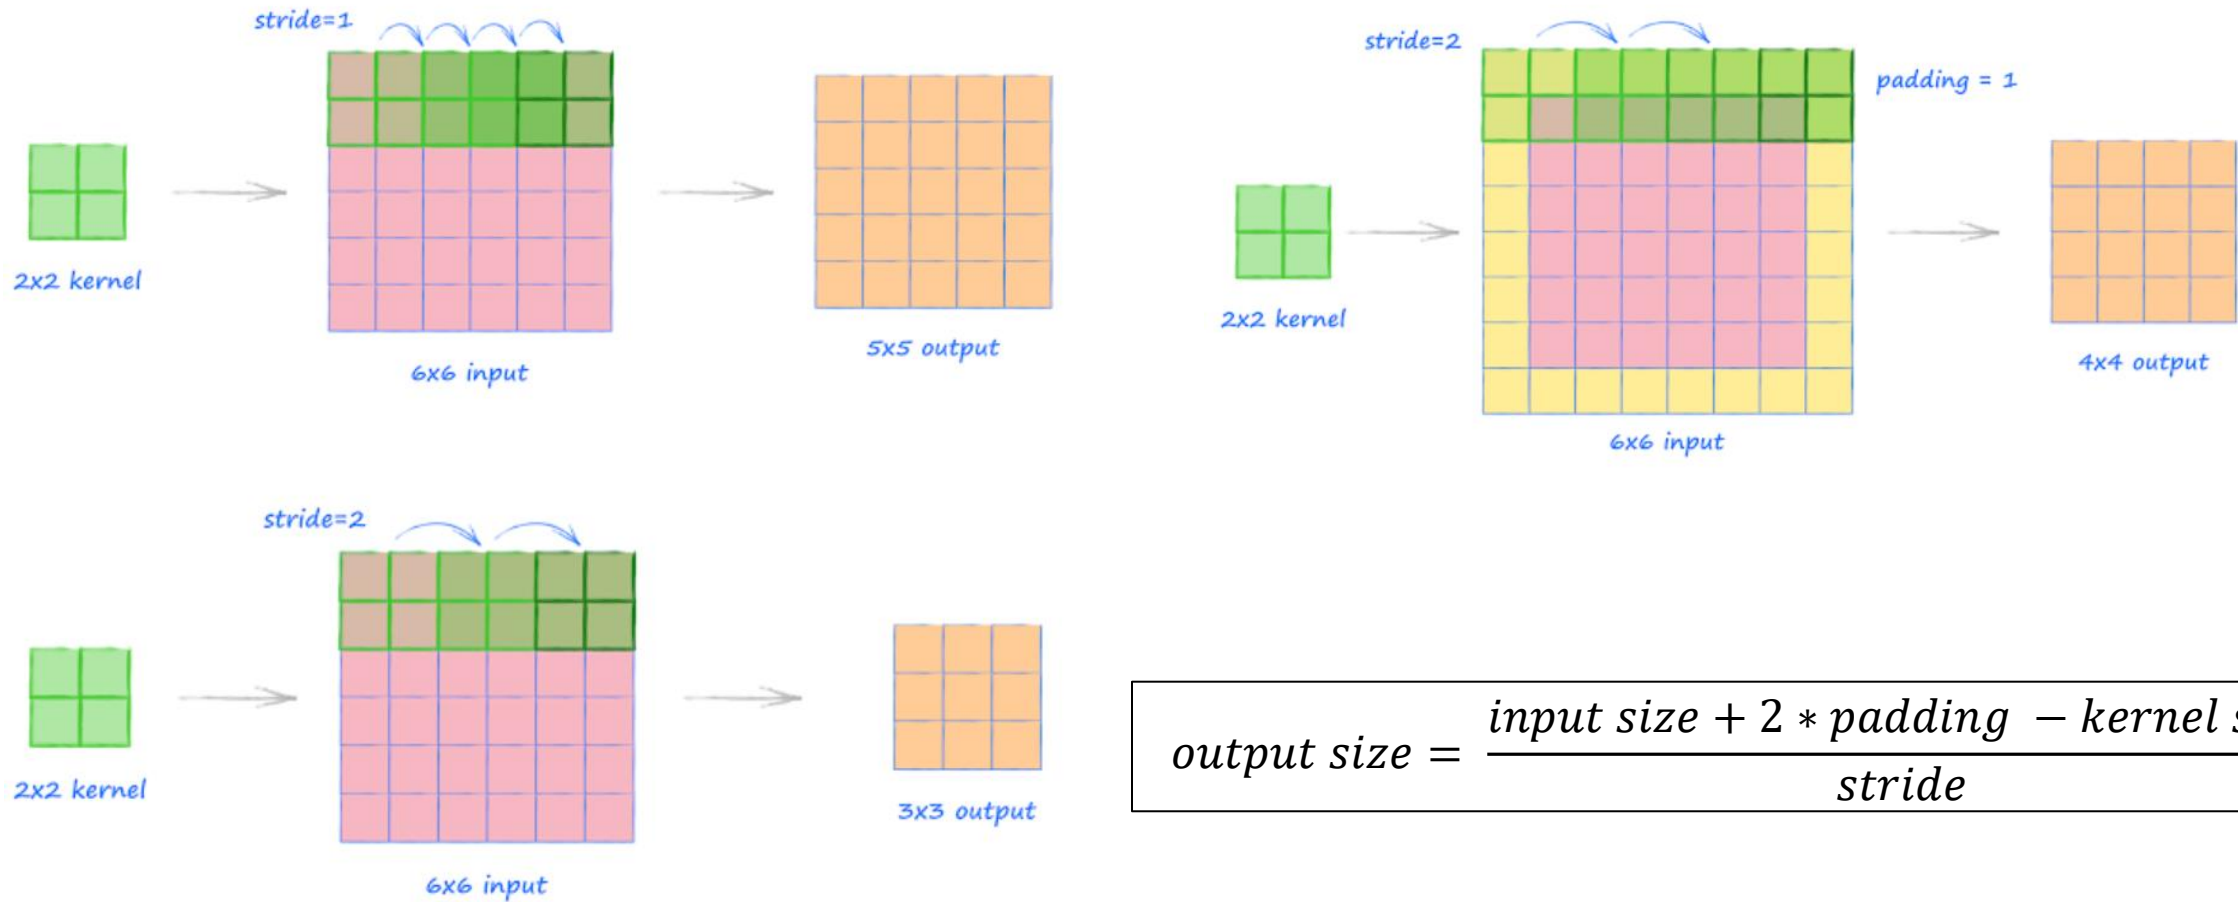
DCGAN

Discriminator: Convolution ~ downsampling



DCGAN


Discriminator: Convolution ~ downsampling



$$\text{output size} = \frac{\text{input size} + 2 * \text{padding} - \text{kernel size}}{\text{stride}} + 1$$

DCGAN

Generator: transposed convolution ~ upsampling




4 step process:

1. Calculate new parameters z and p'
 $z = s - 1; p' = k - p - 1$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**

DCGAN

Generator: transposed convolution ~ upsampling

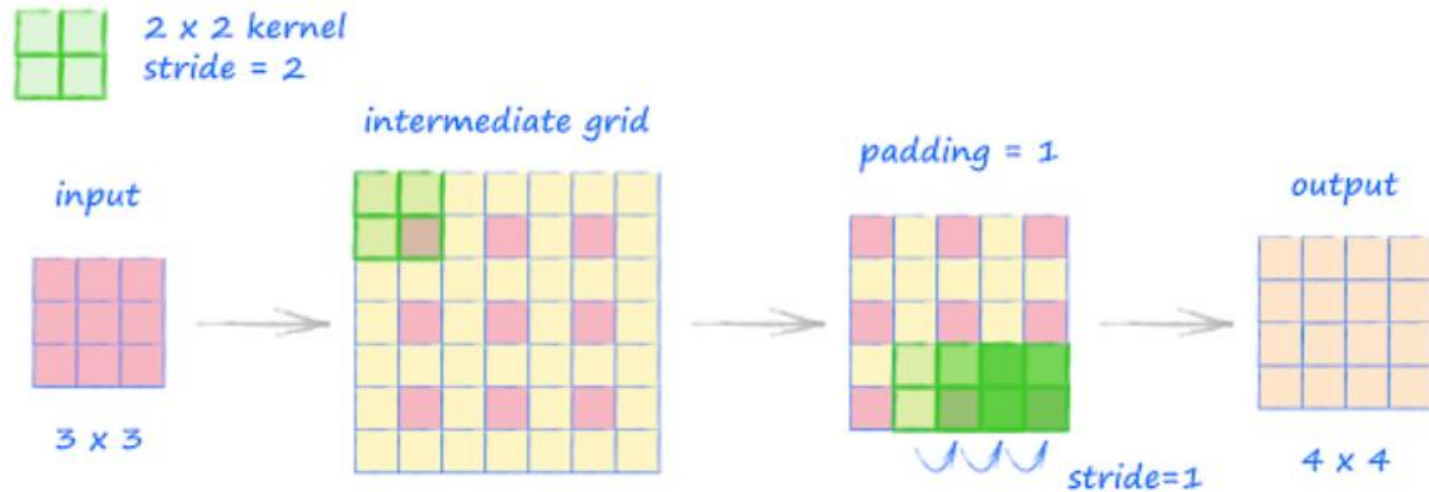


4 step process:

1. Calculate new parameters z and p'
 $z = s - 1$; $p' = k - p - 1$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**

DCGAN

Generator: transposed convolution ~ upsampling



4 step process:

1. Calculate new parameters z and p'
 $z = s - 1$; $p' = k - p - 1$
2. Insert **z number of zeros** between each rows and columns of the input
3. **Pad** the modified input images with **p' number of zeros**
4. Perform standard convolution with **stride length of 1**


$$\text{output size} = (\text{input size} - 1) * \text{stride} - 2 * \text{padding} + \text{kernel size}$$

EXTRA MATERIALS

Supervised vs. Unsupervised Learning


Supervised

- Data: (x, y)
x is data, y is label
- Goal: Learn function to map $x \rightarrow y$
- Examples: Regression, classification, object detection, etc.



Unsupervised

- Data: x
x is data, no labels
- Goal: Learn some hidden or underlying structure of the data
- Examples: Clustering, dimensionality reduction, feature learning, etc.




DCGAN

Batch Normalization

- **Stabilize generator's learning process**
- **Prevent mode collapse**

10 modes


Collapse to 1 mode



DCGAN

Activations:


ReLU



$$a = \max(0, z)$$

**Every layer of generator,
except the last layer**


tanh



$$a = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Last layer of generator

Leaky ReLU



$$a = \max(0.2z, z)$$

Every layer of discriminator